

Math 660: Problem Set 6

Matthew Grasinger

May 3, 2017

1 C1: Linear Iterative Methods

1.1 Source Code

```

1  include("hw6_helpers.jl");
2
3  function jacobi(h::Real; tol::Real=1e-7, max_iterations::Int=10000)
4      grid = Grid(h);
5      enforce_bcs!(grid);
6      next_grid = Grid(grid);
7      const start_time = time();
8      for iteration = 1:max_iterations
9          for j=2:grid.ny-1, i=2:grid.nx-1
10             next_grid.us[i, j] = (grid.us[i-1, j] + grid.us[i+1, j] +
11                                   grid.us[i, j-1] + grid.us[i, j+1] -
12                                   h*h*force(grid.xs[i, j], grid.ys[i, j])) / 4.0;
13          end
14          if norm(next_grid.us - grid.us, 2) < tol
15              return next_grid, iteration, time() - start_time, true;
16          end
17          copy_grid_values!(grid, next_grid);
18      end
19      return grid, max_iterations, time() - start_time, false;
20  end
21
22  function gauss_seidel(h::Real; tol::Real=1e-7, max_iterations::Int=10000)
23      grid = Grid(h);
24      enforce_bcs!(grid);
25      next_grid = Grid(grid);
26      const start_time = time();
27      for iteration = 1:max_iterations
28          for j=2:grid.ny-1, i=2:grid.nx-1
29              # use updated neighboring values
30              next_grid.us[i, j] = (next_grid.us[i-1, j] + next_grid.us[i+1, j] +
31                                    next_grid.us[i, j-1] + next_grid.us[i, j+1] -
32                                    h*h*force(grid.xs[i, j], grid.ys[i, j])) / 4.0;
33          end
34          if norm(next_grid.us - grid.us, 2) < tol
35              return next_grid, iteration, time() - start_time, true;
36          end
37          copy_grid_values!(grid, next_grid);
38      end
39      return grid, max_iterations, time() - start_time, false;
40  end
41
42  function SOR(h::Real; tol::Real=1e-7, max_iterations::Int=10000)
43      const omega = 2.0 / (1.0 + pi * h);
44      grid = Grid(h);
45      enforce_bcs!(grid);
46      next_grid = Grid(grid);
47      const start_time = time();
48      for iteration = 1:max_iterations
49          for j=2:grid.ny-1, i=2:grid.nx-1
50              # use updated neighboring values
51              next_grid.us[i, j] += omega * (
52                  (next_grid.us[i-1, j] + next_grid.us[i+1, j] +
53                    next_grid.us[i, j-1] + next_grid.us[i, j+1] -
54                    h*h*force(grid.xs[i, j], grid.ys[i, j])) / 4.0 -
55                  next_grid.us[i, j]);
56          end
57          if norm(next_grid.us - grid.us, 2) < tol
58              return next_grid, iteration, time() - start_time, true;
59          end
60          copy_grid_values!(grid, next_grid);
61      end
62      return grid, max_iterations, time() - start_time, false;
63  end
64

```

```

65 for (method_str, approx_method) in [("Jacobi iterative solution", jacobi);
66                                     ("Gauss-Seidel iterative solution",
gauss_seidel);
67                                     ("SOR iterative solution", SOR)]
68     for h in [0.1; 0.05; 0.025]
69         println(method_str, ", h = ", h);
70         grid_approx_soln, iterations, time_elapsed, did_converge = approx_method(h);
71         println("    time elapsed:      ", time_elapsed);
72         println("    iterations:          ", iterations);
73         println("    sec/iteration:        ", time_elapsed / iterations);
74         println("    converged:           ", (did_converge) ? "true" : "false");
75         grid_analytical_soln = Grid(h, analytical_soln);
76         println("    rel. L2 error:       ",
77               norm(grid_analytical_soln.us - grid_approx_soln.us, 2) /
78               norm(grid_analytical_soln.us, 2));
79         println();
80     end
81 end

```

1.2 Results

Jacobi iterative solution, $h = 0.1$

time elapsed: 0.004197835922241211
iterations: 286
sec/iteration: 1.4677747979864373e-5
converged: true
rel. L2 error: 5.604084855203742e-5

Jacobi iterative solution, $h = 0.05$

time elapsed: 0.08460402488708496
iterations: 1098
sec/iteration: 7.70528459809517e-5
converged: true
rel. L2 error: 1.4088944068211665e-5

Jacobi iterative solution, $h = 0.025$

time elapsed: 0.682297945022583
iterations: 4180
sec/iteration: 0.0001632291734503787
converged: true
rel. L2 error: 2.14377570070329e-6

Gauss-Seidel iterative solution, $h = 0.1$

time elapsed: 0.0020170211791992188
iterations: 152
sec/iteration: 1.3269876178942228e-5
converged: true
rel. L2 error: 5.624218515105956e-5

Gauss-Seidel iterative solution, $h = 0.05$

time elapsed: 0.0244901180267334
iterations: 580
sec/iteration: 4.222434142540241e-5
converged: true
rel. L2 error: 1.451514372688249e-5

Gauss-Seidel iterative solution, $h = 0.025$

time elapsed: 0.3040142059326172
iterations: 2207
sec/iteration: 0.0001377499800329031
converged: true
rel. L2 error: 2.985097232356857e-6

SOR iterative solution, $h = 0.1$

time elapsed: 0.0005340576171875
iterations: 37
sec/iteration: 1.4433989653716216e-5
converged: true
rel. L2 error: 5.639472054373996e-5

```
SOR iterative solution, h = 0.05
time elapsed:      0.0026369094848632812
iterations:        71
sec/iteration:     3.713957020934199e-5
converged:         true
rel. L2 error:     1.49024396571672e-5
```

```
SOR iterative solution, h = 0.025
time elapsed:      0.017843961715698242
iterations:        139
sec/iteration:     0.0001283738252927931
converged:         true
rel. L2 error:     3.8099858389904924e-6
```

1.3 Discussion

The time per iteration was consistently on the order of 10^{-4} – 10^{-5} seconds per iteration for each method, with longer times corresponding to finer grids. The fact that the time per iteration is consistent across methods agrees with theory. The accuracy, with respect to the analytical solution, is also consistent across methods and appears to depend primarily on the fineness of the grid. Therefore, the most apparent difference between methods is the number of iterations required to reach convergence. The Jacobi method required the largest number of iterations to converge. Gauss-Seidel converged in about half as many iterations on average. SOR converged consistently in the least number of steps, requiring an order of magnitude less steps than either the Jacobi or Gauss-Seidel methods.

2 C2: Conjugate Gradient

2.1 Source Code

```

1  include("hw6_helpers.jl");
2
3  function conjugate_gradient(h::Real; tol::Real=1e-7)
4
5      # initialize data
6      grid = Grid(h);
7      enforce_bcs!(grid);
8      us = grid.us;
9      rs = zeros(grid.nx, grid.ny);
10     qs = zeros(grid.nx, grid.ny);
11     const start_time = time();
12
13     for j = 2:grid.ny-1, i = 2:grid.nx-1
14         x, y = grid.xs[i, j], grid.ys[i, j];
15         rs[i, j] = (-h*h * force(x, y) + us[i+1, j] + us[i-1, j] + us[i, j+1] +
16             us[i, j-1] - 4.0 * us[i, j]);
17     end
18     r2_prev = dot(rs, rs);
19     ps = copy(rs);
20     for j = 2:grid.ny-1, i = 2:grid.nx-1
21         qs[i, j] = 4*rs[i, j] - rs[i+1, j] - rs[i-1, j] - rs[i, j+1] - rs[i, j-1];
22     end
23
24     pdotq = dot(ps, qs);
25     for iteration = 1:(grid.nx*grid.ny)
26         alpha = r2_prev / pdotq;
27         for j = 2:grid.ny-1, i = 2:grid.nx-1
28             us[i, j] += alpha * ps[i, j];
29         end
30
31         if norm(alpha * ps, 2) < tol
32             return us, iteration, time() - start_time, true;
33         end
34
35         for j = 2:grid.ny-1, i = 2:grid.nx-1
36             rs[i, j] -= alpha * qs[i, j];
37         end
38         r2_new = dot(rs, rs);
39         beta = r2_new / r2_prev;
40         for j = 2:grid.ny-1, i = 2:grid.nx-1
41             ps[i, j] = rs[i, j] + beta * ps[i, j];
42             qs[i, j] = (4*rs[i, j] - rs[i+1, j] - rs[i-1, j] - rs[i, j+1] -
43                 rs[i, j-1] + beta * qs[i, j]);
44         end
45         r2_prev = r2_new;
46         pdotq = dot(ps, qs);
47     end
48
49     return us, grid.nx*grid.ny, time() - start_time, false;
50 end
51
52 for h in [0.1; 0.05; 0.025]
53     println("Conjugate gradient, h = ", h);
54     approx_soln, iterations, time_elapsed, did_converge = conjugate_gradient(h);
55     println("    time elapsed:      ", time_elapsed);
56     println("    iterations:        ", iterations);
57     println("    sec/iteration:      ", time_elapsed / iterations);
58     println("    converged:         ", (did_converge) ? "true" : "false");
59
60     const n = size(approx_soln, 1);
61     num = 0.0;
62     den = 0.0;
63     for (j, y) in zip(1:n, 0.0:h:1.0), (i, x) in zip(1:n, 0.0:h:1.0)
64         num += (approx_soln[i, j] - analytical_soln(x, y))^2;

```

```
65     den += (analytical_soln(x, y))^2;
66 end
67
68 println("    rel. L2 error:    ", sqrt(num / den));
69 println();
70 end
```

2.2 Results

```
Conjugate gradient, h = 0.1
  time elapsed:      0.0004241466522216797
  iterations:        28
  sec/iteration:      1.5148094722202846e-5
  converged:          true
  rel. L2 error:      5.6416793358370105e-5

Conjugate gradient, h = 0.05
  time elapsed:      0.0019061565399169922
  iterations:        58
  sec/iteration:      3.286476792960331e-5
  converged:          true
  rel. L2 error:      1.4923993244701214e-5

Conjugate gradient, h = 0.025
  time elapsed:      0.013283014297485352
  iterations:        118
  sec/iteration:      0.00011256791777529959
  converged:          true
  rel. L2 error:      3.835920588689894e-6
```

2.3 Discussion

The time per iteration was on the order of 10^{-4} – 10^{-5} seconds per iteration for the conjugate gradient method, which is consistent with the methods used in Section 1. The accuracy, with respect to the analytical solution, was also consistent with the methods used in Section 1. However, the conjugate gradient required less iterations to converge for all cases.